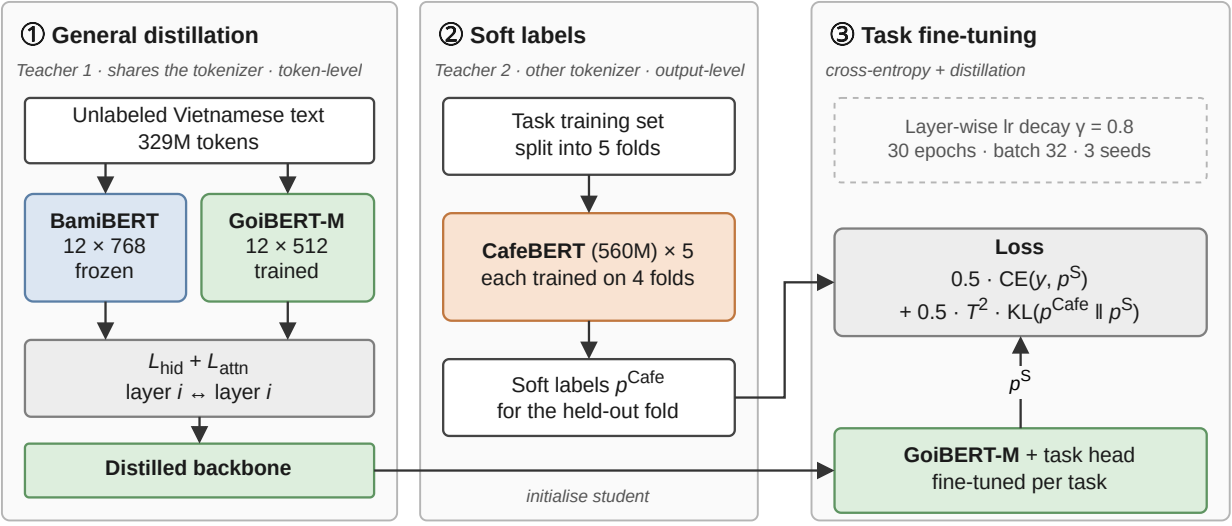
① General distillation
Teacher 1 · shares the tokenizer · token-level
Unlabeled Vietnamese text
329M tokens
BamiBERT
12 × 768
frozen
GoiBERT-M
12 × 512
trained
$L_{hid} + L_{attn}$
layer i ↔ layer i
Distilled backbone
② Soft labels
Teacher 2 · other tokenizer · output-level
Task training set
split into 5 folds
CafeBERT (560M) × 5
each trained on 4 folds
Soft labels p^{Cafe}
for the held-out fold
initialise student
③ Task fine-tuning
cross-entropy + distillation
Layer-wise lr decay γ = 0.8
30 epochs · batch 32 · 3 seeds
Loss
$0.5 \cdot CE(y, p^S)$
$+ 0.5 \cdot T^2 \cdot KL(p^{Cafe} \| p^S)$
p^S
GoiBERT-M + task head
fine-tuned per task